

Distributed Computing (2020)

Synchronization


Christian' Algorithm

Algorithm

Let S be the timeserver and T_s be its time.

- Process P requests the time from S .
- After receiving the request from P , S prepares a response and appends the time T_s from its own clock and sends it back to P
- P then sets its time to be $T_p = T_s + RTT/2$

Figure 7.1(a) illustrates the idea behind the algorithm more clearly.



Berkeley's Algorithm

Algorithm

Elect* the master amongst N nodes. Let T_m be the time estimate of the master's clock.

Let $t[i]$ contain the time at each i slave at master

If master


send its T_m along with query for $t[i]$ to slaves; /* for $i = 1 \dots N-1$ */

Adjust = $\text{Sum}(t[i])/N$ /* take average including masters


send offset $[i] = \text{Adjust} - t[i]$ to each slave;

If slave


sends query response as $t[i] = T_m - T[i]$; /* for $i = 1 \dots N-1$; calculates the difference between master timestamp T_m , and its own timestamp T */



(a) Query-poll



(b) Response



(c) Adjust

Lamport's Logical(scalar) Clock





Figure 7.4 (a) Events a and b are events of process P_1 . (b) Events a and b are send and receive events of process P_1 and P_2 , respectively, of the same message m .

If a and b are two events, $a \rightarrow b$ would mean a happened before b :

1. If a and b are events internal in the same process.
2. If a is the event corresponding to the sending of message m in one process, and b is the event corresponding to receiving of the same message m in the other process.

Lamport's Logical Clock



Vector Timestamp Ordering

Each process P_i maintains a vector of integer clock with elements of the vector being N , N is the number of processes.

Algorithm

For every local event,

- $V_i[i] = V_i[i] + 1;$ /* increment only i th element in the vector


For every send event at P_i ,

- $V_i[i] = V_i[i] + 1;$ /* increment only i th element in the vector
- Send the Message + V_i to the receiver;

For every received message at P_i from P_j

- $V_i[i] = V_i[i] + 1;$
- $V_i[j] = \max(V_i[j], V_j[j]);$ /* for $j \neq i$

Vector Timestamp Ordering



Causal events:

$H \rightarrow G : [0,0,1] < [2,3,1]$ /* any one of the element is smaller*/

$F \rightarrow L : [2,2,1] < [2,4,4]$

$A \rightarrow G : [1,0,0] < [2,3,1]$ /* all elements are smaller*/

Concurrent events:

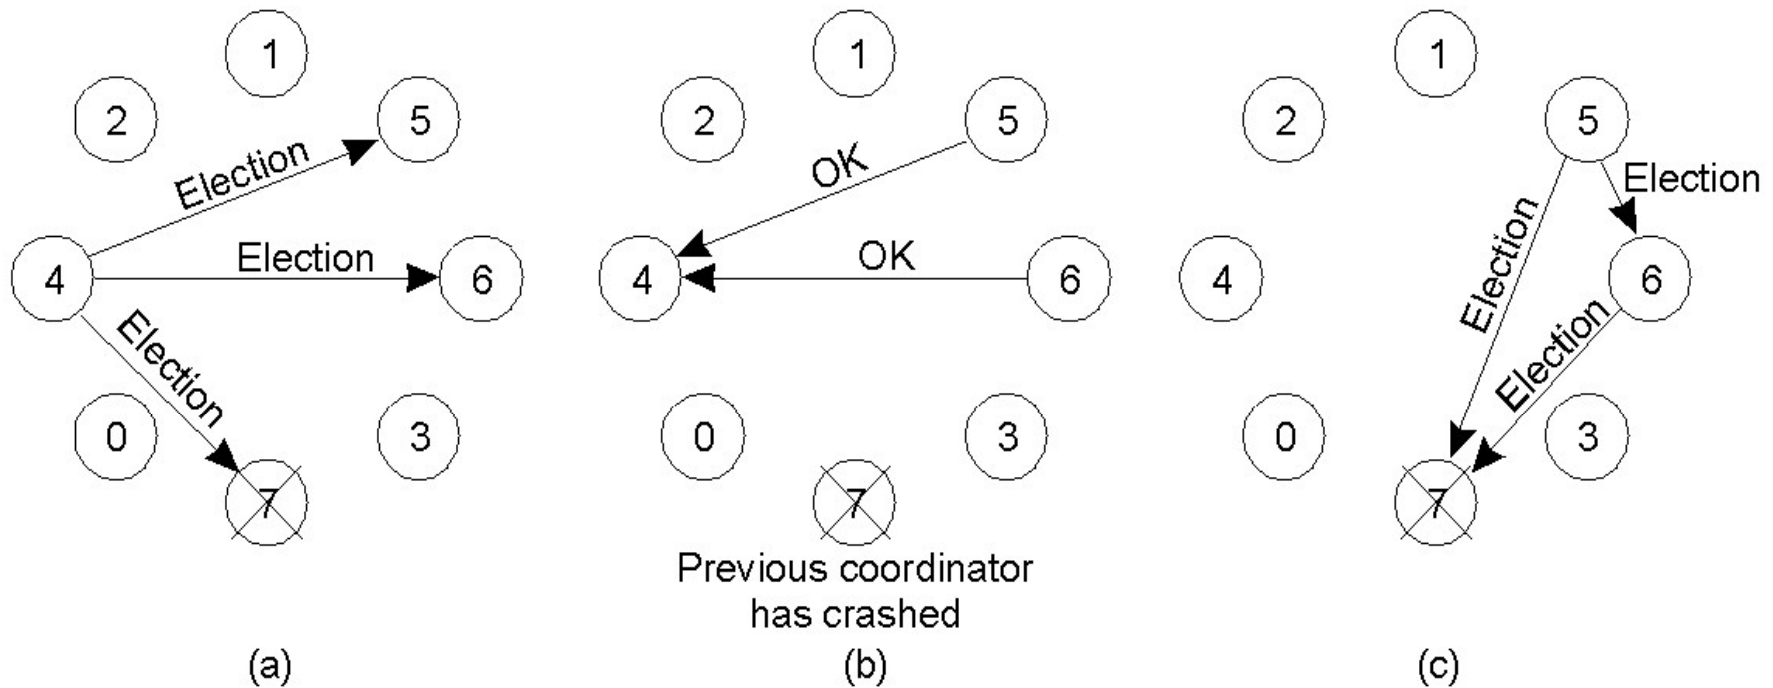
C and F : $[3,0,0] \parallel [2,2,1]$ /* $3 > 2 ; 0 < 1$ */

F and J : $[2,2,1] \parallel [0,0,2]$ /* $2 > 0 ; 1 < 2$ */

I and C : $[0,0,1] \parallel [3,0,0]$ /* $0 < 3 ; 1 > 0$ */

Coordinator Selection Election Algorithms



The Bully Algorithm (1)



The bully election algorithm

- Process 4 holds an election
- Process 5 and 6 respond, telling 4 to stop
- Now 5 and 6 each hold an election

Global State (3)



- d) Process 6 tells 5 to stop
- e) Process 6 wins and tells everyone

Algorithm

1. Start of ELECTION

- Process P_i initiates an election if it just recovered from failure or it notices that the coordinator has failed.
 - P_i sends ELECTION(i) messages to all processes P_j with higher IDs ($j > i$)
 - Awaits OK messages.
 - If timeout */*declares itself as the coordinator process.**
 - Sends COORDINATOR(i) message to all lower ID nodes P_j ($j < i$)
- /* As no higher ID process responded, it is the highest ID process by default */*

2. P_j receives ELECTION(i) message

- If $j > i$, send OK to P_i
- Send ELECTION(j) messages to higher ID process.

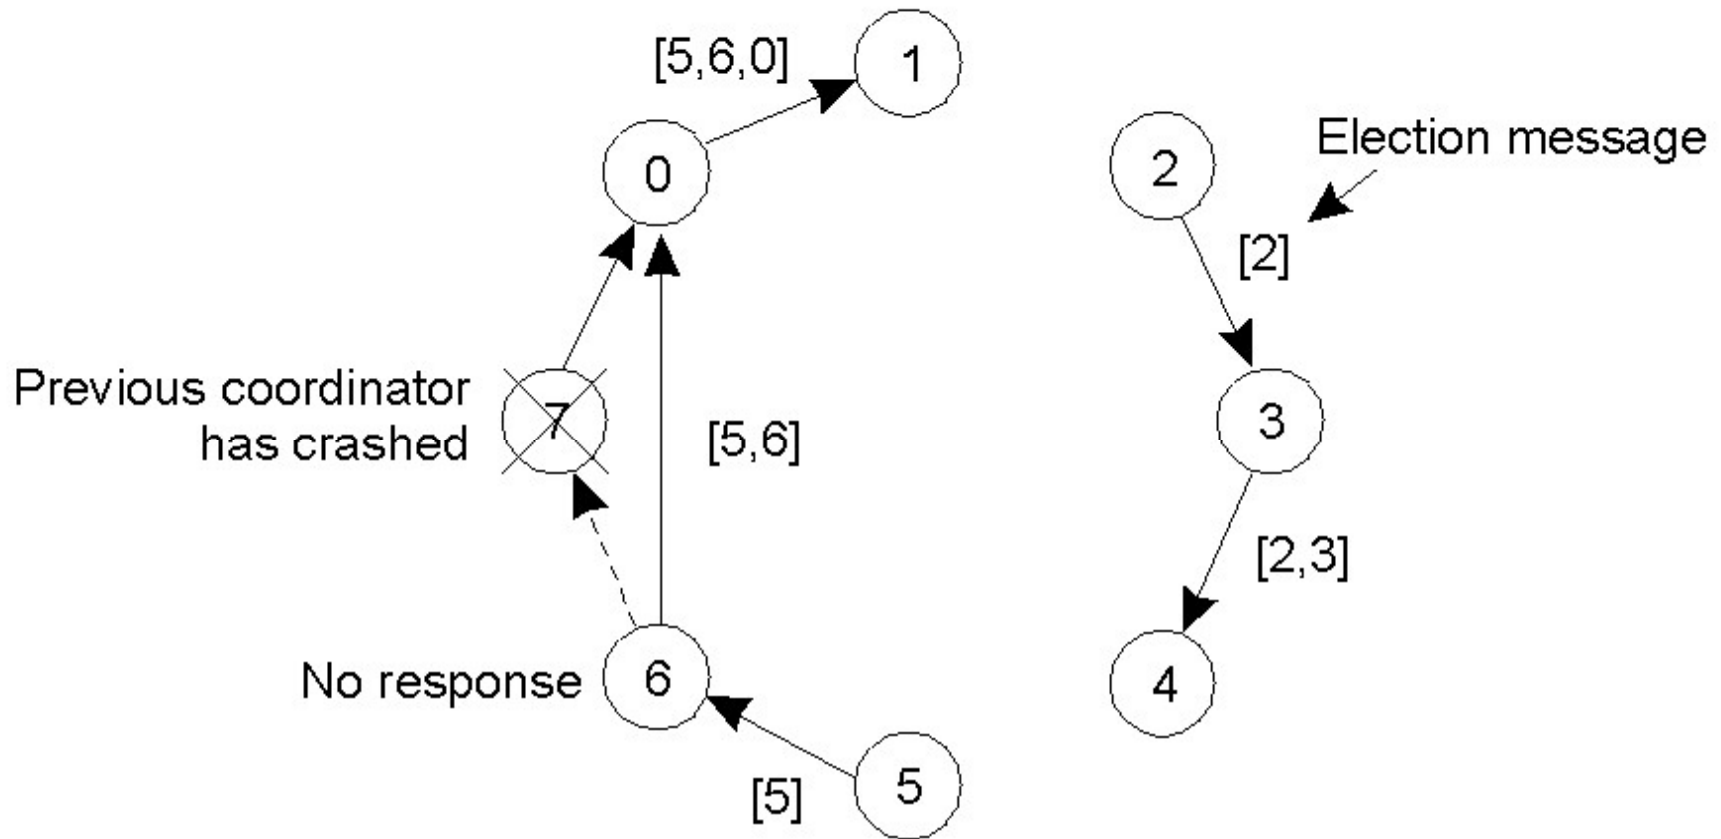
3. P_i receives OK from P_j

- Stop ELECTION(i) process, */*higher ID process has been found*/*
 - Awaits, COORDINATOR(j) message from any P_j ($j > i$)
 - Else sends COORDINATOR(i) message to all lower ID nodes P_j ($j < i$)
- /* because no higher ID process responded, so it is the highest ID process*/*
- /* Some versions of algorithm suggest to restart the ELECTION process instead */*

4. P_i receives COORDINATOR(j) from P_j

- It stops its election process and P_j is termed as elected coordinator node

A Ring Algorithm



Election algorithm using a ring.

Algorithm

1. Start of ELECTION

- P_i initiates an election if it just recovered from failure or it notices that the coordinator has failed.
- The ELECTION[i] message along with the identifier of the node is sent to the next downstream process alive and forwarded along the ring by each process.

2. Receiving an ELECTION message

- If P_i receives the its own ELECTION message, it
 - Removes the message and selects the highest ID process from the message.
 - Sends a COORDINATOR message with ID of the highest process
- Else, if P_j receives the process, i.e., ($j \neq i$), it
 - Appends its own ID into the message.
 - Forwards the ELECTION [] message to the next downstream node.

3. Receiving COORDINATOR message

- If P_i receives the COORDINATOR message, it
 - Removes this message from the ring. /* the message has traversed the ring*/
 - Marks the election process to be over.
- For other processes P_j ($j \neq i$)
 - The ID of the coordinator is noted.
 - Forward the message to downstream process.

Requirements of Mutual Exclusion Algorithms

1. Safety property: At most one process may execute in the critical region (CR) at a time.
2. Liveness property: A process requesting entry to the CR is eventually granted it. There should not be deadlock and starvation
3. Fairness: Each process should get a fair chance to execute the CR.


Performance Metrics

1. Synchronization delay: Time interval between critical region (CR) exit and new entry by any process.
2. System throughput: Rate at which requests for the CR get executed.
3. Message complexity: Number of messages that are required per CR execution by a process.
4. Response time: Time interval from a request send to its CR execution completed.


Classification of Mutual Exclusion Algorithms

- Centralized mutual exclusion algorithms
- Distributed mutual exclusion algorithms
- Non-token-based algorithm:
 - Lamport's Distributed Mutual Algorithm
 - Ricart–Agrawala Algorithm
 - Maekawa's Algorithm
- Token-based algorithm:
 - Suzuki–Kasami's Broadcast Algorithm
 - Singhal's Heuristic Algorithm
 - Raymond's Tree-Based Algorithm


Mutual Exclusion: A Centralized Algorithm



(a)



(b)



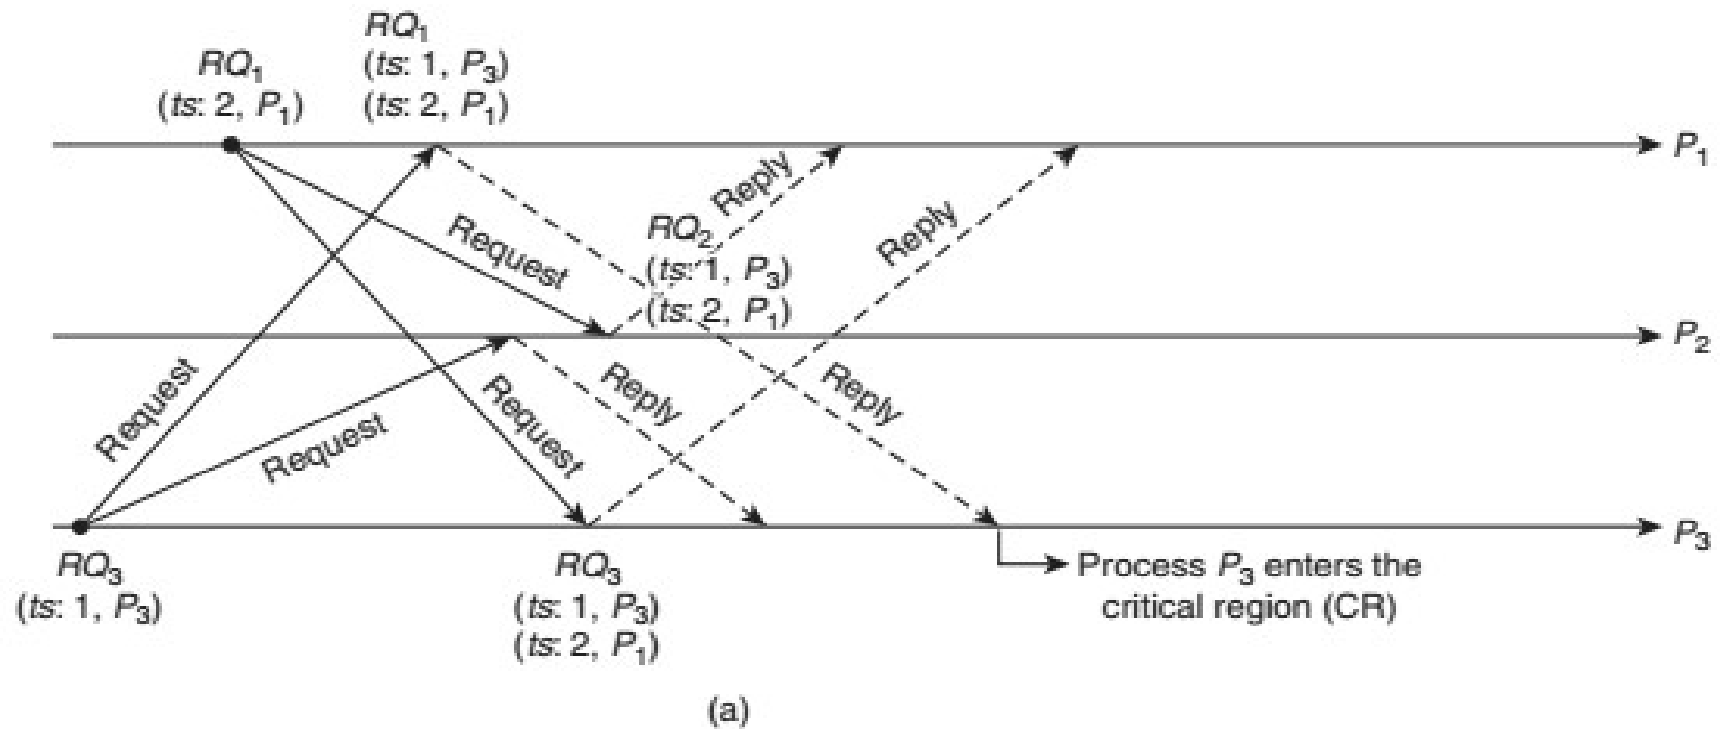
(c)


- a) Process 1 asks the coordinator for permission to enter a critical region. Permission is granted
- b) Process 2 then asks permission to enter the same critical region. The coordinator does not reply.
- c) When process 1 exits the critical region, it tells the coordinator, when then replies to 2

Performance Parameters

1. The algorithm is simple and fair as it handles the request in the sequential order.
2. It guarantees no starvation.
3. It uses three messages as REQUEST, REPLY and RELEASE.
4. It has a single point of failure.
5. Coordinator selection could increase synchronization delay especially at times of frequent failures.

Lamport's Distributed Mutual Algorithm





Algorithm


1. The critical region request message from P_i
 - P_i sends a REQUEST(timestamp $_i$, i) to all P_j in its quorum set R_i (for $j \neq i$)
2. P_j on receiving the REQUEST message from P_i
 - If variable voted $_j$ = True or if the process P_j itself is currently executing the CR
 - then REQUEST from P_i is deferred and pushed in the queue RD_j .
 - else REPLY is send to P_i /* termed as permission granted*/
and variable voted $_j$ is set to TRUE
3. The execution of the CR
 - If all REPLY received from processes in R_i , enter CR.
4. The release of the CR: P_i
After execution of CR,
Process P_i sends RELEASE to all P_j in R_i .

5. The receipt of RELEASE message: $/* P_i$ sends to all P_j in R_i^* /
- If RD_j is nonempty
 - (a) dequeue top of the queue RD_j ,
 $/*$ looks out for requests which came while P_i was in CR^* /
 - (b) P_j sends a REPLY message to only this dequeued process.
 - (c) Set $voted_j$ to be TRUE
 - If queue RD_j was empty
 - (a) $voted_j = \text{false}$

Performance Parameters

- Lamport's algorithm has message overhead of total $3(N - 1)$ messages: $N - 1$ REQUEST messages to All process (N minus itself), $N - 1$ REPLY messages, and $N - 1$ RELEASE messages per CR invocation.
- The synchronization delay is T . Throughput is $1/(T + E)$.
- The algorithm has been proven to be fair and correct. It can also be optimized by reducing the number of RELEASE messages sent.

Ricart-Agrawala Algorithm



Performance Parameters


The algorithm does not use explicit RELEASE message. The dequeuing is done on the receipt of REPLY itself. Thus, total message overhead would be $2(N - 1)$ messages, that is, for entering a CR, $(N - 1)$ requests and exiting $(N - 1)$ replies.

2. The failure of any process almost halts the algorithm (recovery measures are needed) as it requires all replies.

Maekawa's Algorithm

Properties of a Quorum Set

1. $(R_i \cap R_j \neq \text{Null})$, (for all i and j in $i \neq j, 1 \leq i, j \leq N$)
/*Intersection of any two Quorum sets should not be Null*/
2. $(P_i \in R_i)$, (for all $i, 1 \leq i \leq N$)
/*Each process belongs to its own quorum set R_i */
3. $(|R_i| = K)$, (for all i in $1 \leq i \leq N$)
/*Size of each quorum set is K */
- 4 Any process P_j is contained in some M number of R_i s (for $1 \leq i, j \leq N$.)



(a)


$N = 4, R = 4, K = 2$

$R_1 = \{P_1, P_2\}$

$R_2 = \{P_2, P_3\}$

$R_3 = \{P_3, P_4\}$

$R_4 = \{P_4, P_1\}$




$N = 4, R = 4, K = 3$

$R_1 = \{P_1, P_2, P_4\}$

$R_2 = \{P_2, P_1, P_3\}$

$R_3 = \{P_3, P_2, P_4\}$

$R_4 = \{P_4, P_1, P_3\}$



(c)

$N = 7, R = 7, K = 3$

$R_1 = \{P_1, P_2, P_3\}$

$R_2 = \{P_2, P_4, P_6\}$

$R_3 = \{P_3, P_5, P_7\}$

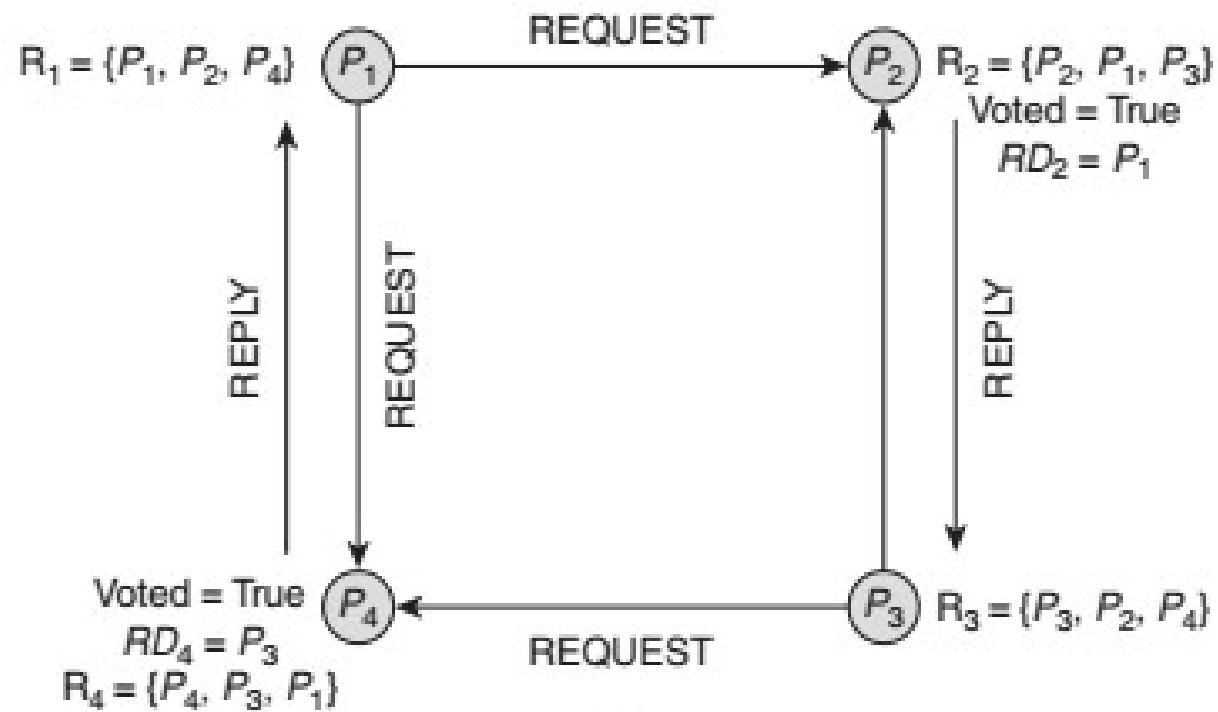
$R_4 = \{P_1, P_4, P_5\}$

$R_5 = \{P_5, P_2, P_7\}$

$R_6 = \{P_6, P_1, P_7\}$

$R_7 = \{P_7, P_3, P_4\}$

N = No. of processes; R = No. of Quorums; K = size of Quorum



(d)

Algorithm

1. The critical region request message from P_i
 - P_i sends a REQUEST(timestamp _{i} , i) to all P_j in its quorum set R_i (for $j \neq i$)
2. P_j on receiving the REQUEST message from P_i
 - If variable Voted _{j} = True or if the process P_j itself is currently executing the CR
 - then REQUEST from P_i is deferred and pushed in the queue RD_j .
 - else REPLY is send to P_i /* termed as permission granted*/
and variable Voted _{j} is set to TRUE
3. The execution of the CR
 - If all REPLY received from processes in R_i , enter CR.
4. The release of the CR: P_i
After execution of CR,
Process P_i sends RELEASE to all P_j in R_i .

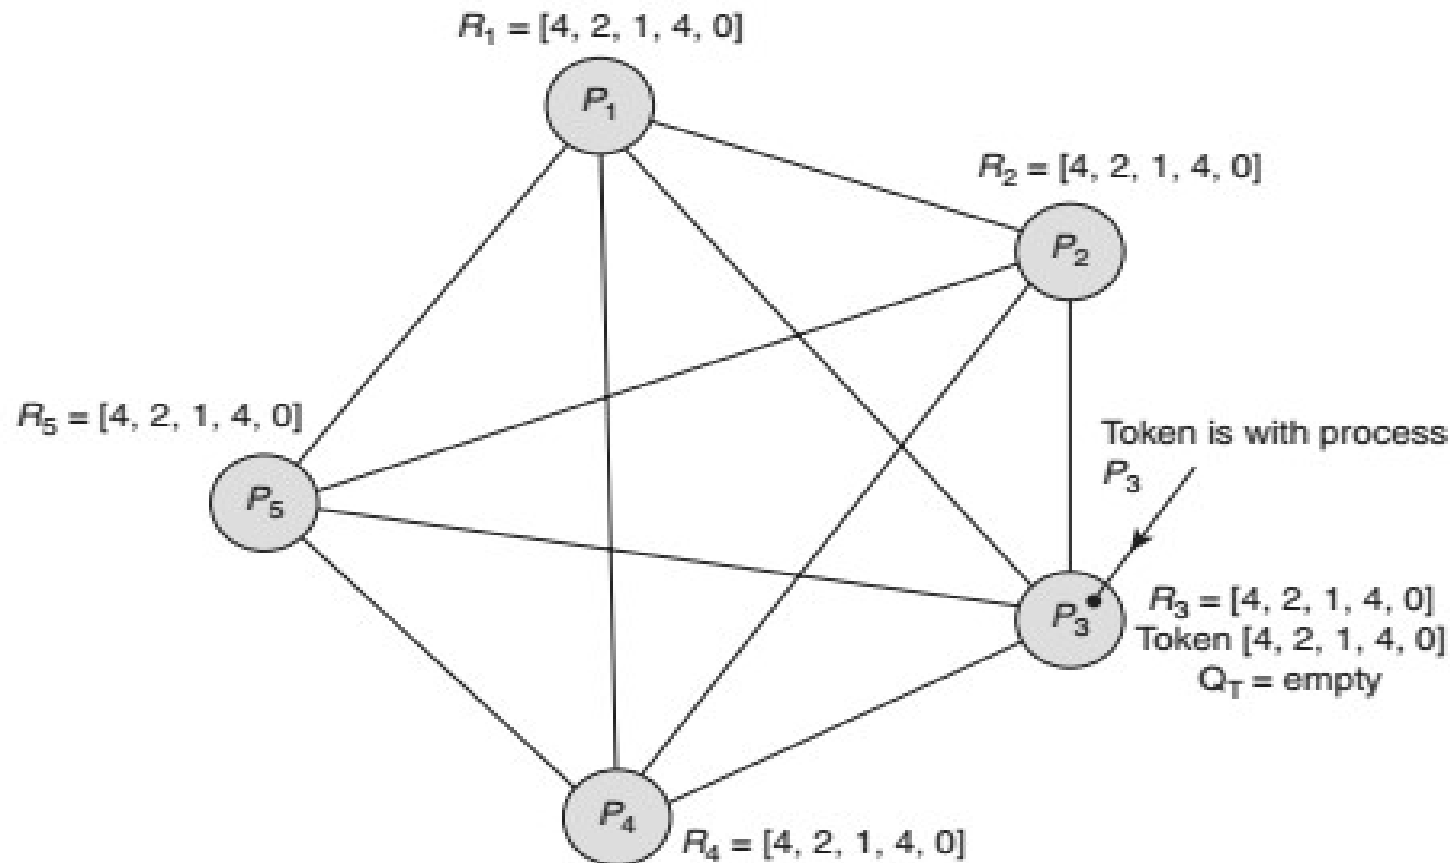
5. The receipt of RELEASE message: */* P_i sends to all P_j in R_i^* */*
 - If RD_j is nonempty
 - dequeue top of the queue RD_j ,
/ looks out for requests which came while P_i was in CR */*
 - P_j sends a REPLY message to only this dequeued process.
 - Set $Voted_j$ to be TRUE
 - If queue RD_j was empty
 - $Voted_j = \text{false}$

Performance Parameters


1. An execution of the CR requires N REQUEST, N REPLY and N RELEASE messages, thus requiring total $3N$ messages per CR execution.
2. Synchronization delay is $2T$.
3. $M = K = N$ works best.

Token-Based Algorithms

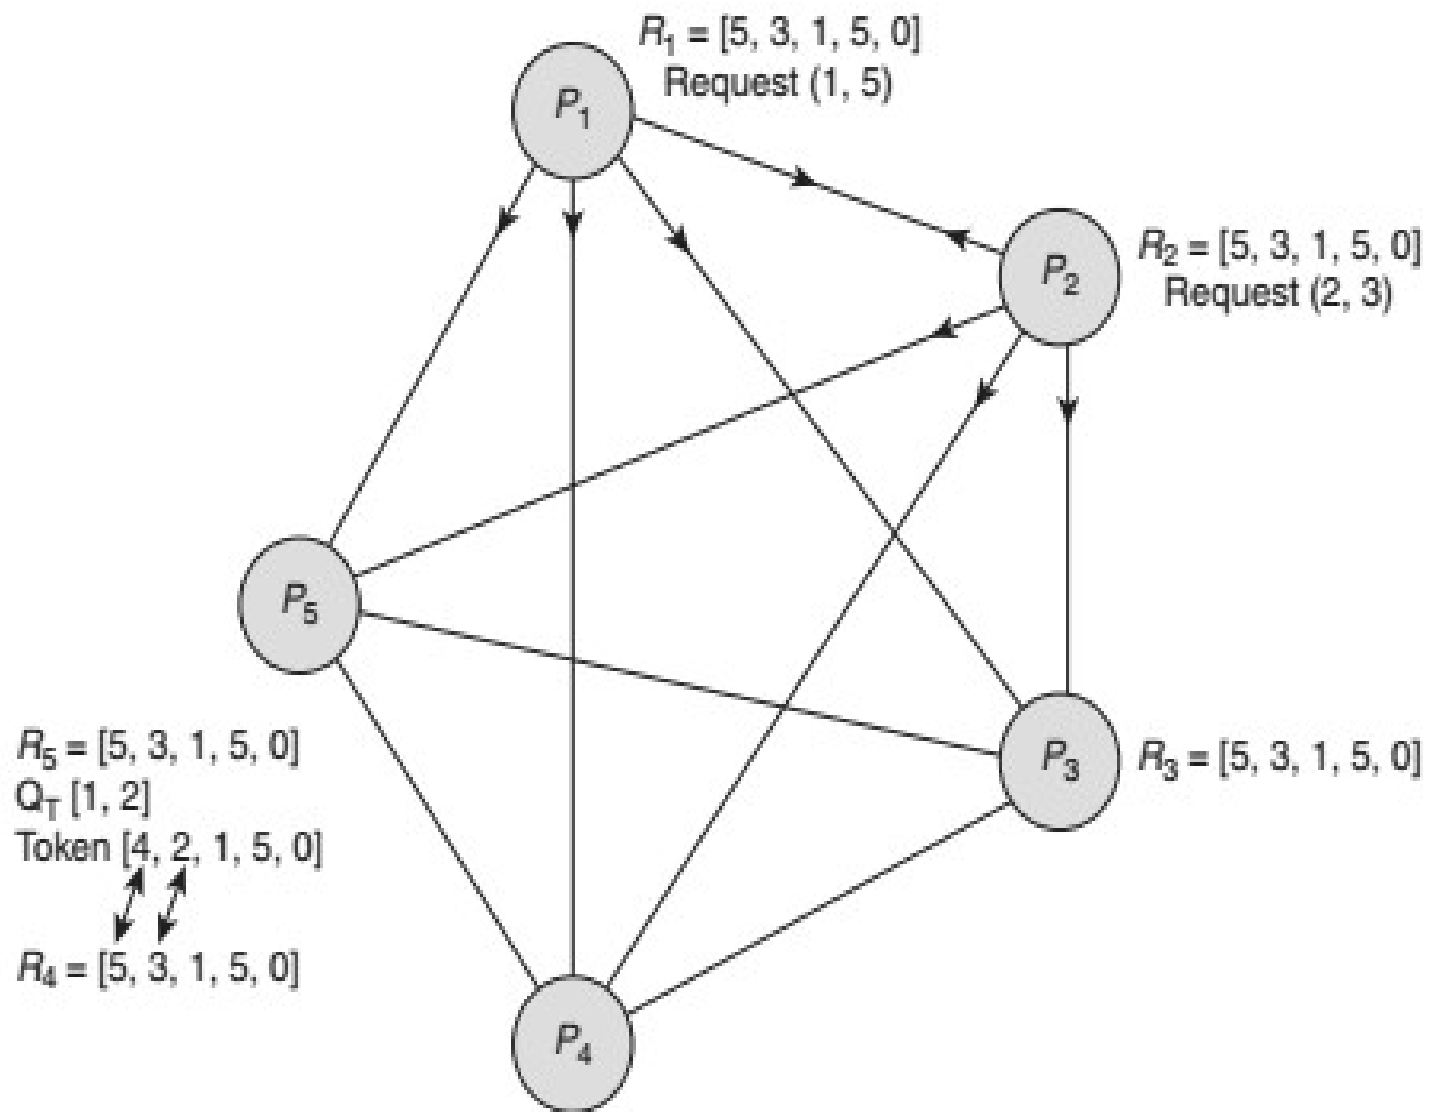
Suzuki–Kasami's Broadcast Algorithm



(a) Initial state: P_3 has the token. No other process is regnesting.
Process P_1, P_2 , and P_4 had earlier executed CR.



(b) Process P_4 wants to enter CR by broadcasting request (4,5)



Algorithm

1. The critical region request message (P_i does not possess the token)
 $n = R_i[i] = R_i[i] + 1$;
REQUEST(i, n) message to all other processes, wait for the token to arrive.
2. P_j on receiving the request message from P_i : REQUEST(i, n)
 - P_j sets $R_j[i] = \max(R_j[i], n)$. /* to take of the outdated requests */
 - If P_j was not executing CR itself and is holding the token
and if $R_j[i] = \text{Token}[i] + 1$;
then it sends the token to P_i
/* token moves to P_i along with its data structure, $\text{Token}[]$ and Q_T */
3. On the execution of the CR: P_i
 - Process P_i executes the CR after it has received the token.
4. On the release of the CR: P_i
 - $\text{Token}[i] = RN_i[i]$ /* update the token */
 - if ($RN_i[j] = \text{Token}[j] + 1$), for all j , puts the ID of P in the Q_T , if not already there.
 - If queue is nonempty
Dequeue the top of the Q_T entry, and pass the token to the process at the top of the queue,
 Q_T

Performance Parameters

1. It is simple and efficient.
2. The algorithm requires at most N messages to obtain the token to enter CR.
3. The synchronization delay in this algorithm is 0 or T . Zero synchronization delay, if the process already holds the token.

Singhal's Heuristic Algorithm

1. E = Executing and having the token
2. H = Holding the token and not executing
3. R = Requesting the token
4. N = Neutral, none of the above

Data Structures

Let there be N processes. A process P_i has following two arrays:

1. Array holding the state of the process $State_i[1...N]$
2. Array $R_i[1....N]$ maintaining the highest number of requests received from each process.

The token has the following two arrays:

1. $T_State[1...N]$, maintaining the states of the processes.
2. $T_R[1...N]$, number of CRs executed by each process.

Arrays Initialization

For each process, P_i for $i = N \dots 1$ do

$State_i[j] = N$; for $j = N \dots k$ do; /* All processes $\geq k$ are set to neutral state = N

$State_i[j] = R$; for $j = k - 1 \dots 1$ do /* All processes $< k$ are assigned the state R

Requesting process

$R_i[j] = 0$; for $j = 1 \dots N$ do /* initially there is no request received by any process

 /* Let P_1 hold the token initially */

$State_1[1] = H$ /* P_1 has the state hold and it has the token */

For the token

$T_State[j] = N$; for $j = 1 \dots N$ /* The token is not being used by any, the array is in state N

$T_R[j] = 0$; for $j = 1 \dots N$ /* There is no process that has executed the CR

Algorithm

1. The critical region request message (P_i does not possess the token)
 - Set $State_i[i] = R$; /*requesting*/
 $n = R_i[i] = R_i[i] + 1$;
REQUEST(i, n) message to all other processes P_j for which $State_i[j] = R$
/* This reduces the number of request messages sent as compared to Suzuki-Kasami algorithm. */
2. P_j on receiving the request message from P_i : REQUEST(i, n)
 - P_j sets $R_j[i] = \max(R_j[i], n)$. /* to take care of the outdated requests*/
 - If $State_j[j] = N$, then set $State_j[i] = R$. /* Update the State of P_i in P_j to the requesting state*/
 - If $State_j[j] = R$, if $State_i[j] \neq R$, Set $State_j[i] = R$, send REQUEST($j, R_j[j]$) to P_i
/* If P_j is requesting, let P_i know of this request*/
 - If $State_j[j] = E$, then Set $State_j[i] = R$
 - If $State_j[j] = H$, then Set $State_j[i] = R$, $T_State[i] = R$, $T_R[i] = n$, $State_j[j] = N$
and send the token to process P_i .
3. On the execution of the CR: P_i
 - Once the token is received, Set $State_i[i] = E$;
 - Process P_i executes the CR after it has received the token.
4. On the Release of the CR: P_i
 - Set $State_i[i] = N$ and $T_State[i] = N$,
 - For all $j = 1 \dots N$ do
if $State_i[j] > T_State[j]$ /* according to hierarchical order */

then, $T_State[j] = State_i[j]$; $T_R[j] = R_i[j]$; /*update token information from local information*/

else $State_i[j] = T_State[j]$; $R_i[j] = T_State[j]$ /* update local information from the token information */

- If (for all j ; if $State_i[j] = N$), then

Set $State_i[i] = H$

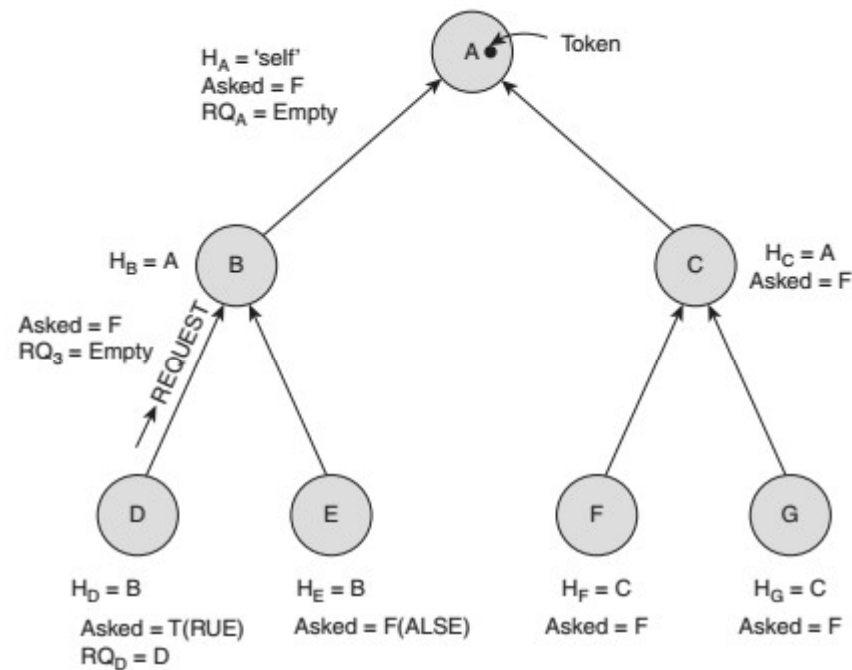
else, if $State_i[j] = R$, send the token to a process P_j

Performance Parameters

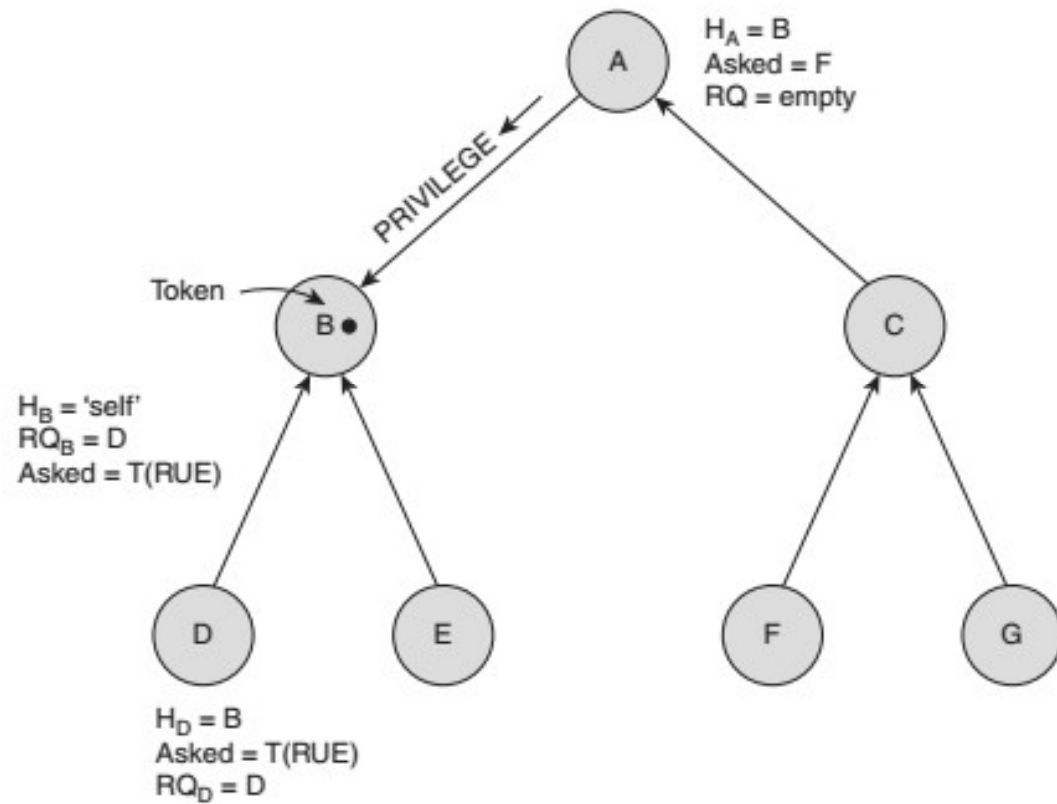
Performance Parameters

1. The number of REQUEST messages can vary from $N/2$ (Average value of the identifier) to N (max value).

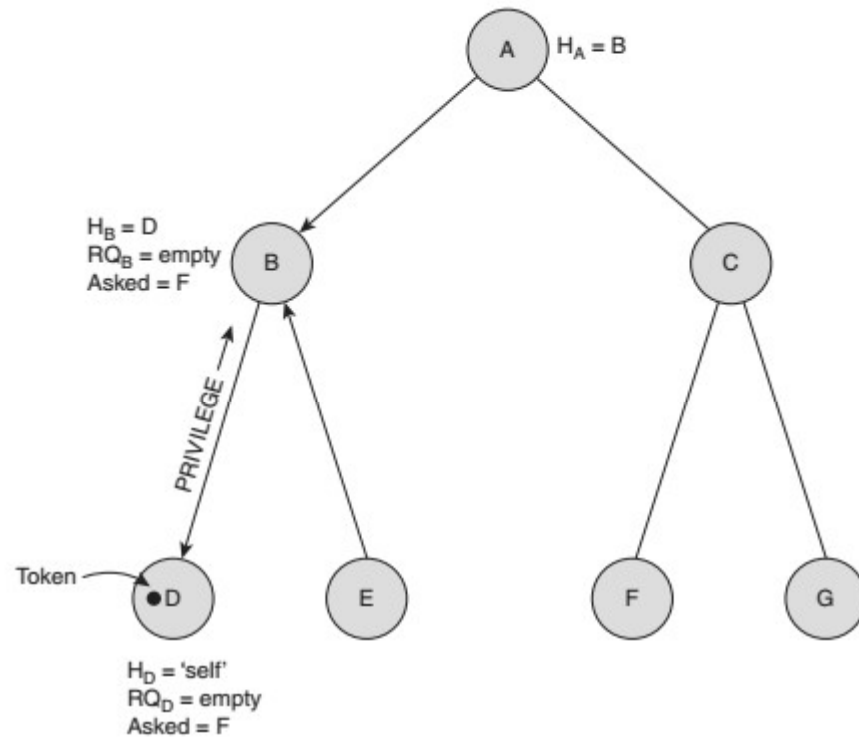
Raymond's Tree-Based Algorithm



(a) Process D sends the REQUEST to process B which forwards it to process A.



(b) Process B receives the token after granting the PRIVILEGE from process A.



(c) Process D receives the token after granting the PRIVILEGE from process B.

Algorithm

1. The critical region request message

- If P_i does not hold the token; /* if token held, no need to send REQUEST*/
and its RQ_i is not empty; /*process requires the token for itself or for its neighbors*/
and Asked = FALSE /*it has not already sent a REQUEST message, either
Then increment its RQ_i
check H_i and send the REQUEST to the process held in variable H_i ;
Asked = TRUE;
else Asked = FALSE;

2. P_j receiving the request message from P_i

- If P_j has the token, i.e., $H_j = \text{'self'}$
and if is not executing the CR and its RQ_j is not empty
and the element at the head of its RQ is not 'self'.
send the PRIVILEGE message to a requesting process, update its $H_j = i$
else, if the process is an unprivileged process
 P_j forwards the request to parent process. /*step 1 is executed then */

3. Receipt of a PRIVILEGE message

- If P_i is the processes which had requested the token, RQ_i has itself at the top and $H_i = \text{'self'}$;

decrement RQ_i 's value for yourself;

Enter the CR; /*Your request was granted*/

Asked = FALSE

- If P_i had forwarded the request on behalf of its neighbor

Dequeue the RQ_j /* or decrement the value*/

send the PRIVILEGE message to a requesting process (top of RQ_j)

update H_i and change the parent

if RQ_i is still not empty,

send the REQUEST message to the new H value

Asked = TRUE;

Else Asked = FALSE;

4. The execution of the CR

- The process P_i gets to execute the CR if the process P_i has the token available only if its own ID is at the head of its RQ;

5. P_i exiting CR

- if RQ_i is nonempty then

Dequeue RQ_i ; Let us call it P_d /* get the process who had send the request*/

send the token to this process P_d ;

change the $H_i = d$; Asked = FALSE;

- if RQ is still not empty, send REQUEST to the parent process, Asked = TRUE;

Performance Parameters

The algorithm exchanges only $O(\log N)$ messages under light load and four messages under heavy load

to execute the CR, where N is the number of nodes in the network.